

# STAT527 Term Project

Miller Kodish, Ian Ou, Vinay Pundith

12/17/2025

## Part 1: Define Helper Functions

- (a) These functions are used for tables. `useful_columns()` checks each column and keeps only the ones that don't have too many missing values, based on a threshold. `df_stats()` then uses that result to report simple summary info about the data, like how many rows it has and how many "useful" columns remain after filtering out columns with lots of NAs.

```
useful_columns <- function(df, na_threshold = 0.85) {
  na_fraction <- sapply(df, function(col) mean(is.na(col)))
  names(na_fraction[na_fraction <= na_threshold])
}

df_stats <- function(df, na_threshold=0.85) {
  cols <- useful_columns(df, na_threshold)
  list(rows = nrow(df), useful_cols = length(cols))
}
```

- (b) These functions are used for plotting. `int_plot()` automatically turns ggplot figures into interactive Plotly plots when the output is HTML, but keeps them static for PDFs. `plot_time_series()` cleans out missing values, makes a simple time-series plot with points and a smooth trend line so we don't have to repeat the same code.

```
# make sure plots are interactive in the R when compiled but static when in PDF
int_plot <- function(p) {
  if (knitr:::is_html_output()) {
    plotly::ggplotly(p)
  } else {
    p
  }
}

plot_time_series <- function(data, y, title, ylab, color_point, color_line) {
  clean_data <- data |> filter(!is.na(testDate), !is.na(.data[[y]]))
  p <- ggplot(clean_data, aes(x = testDate, y = .data[[y]])) +
    geom_point(alpha = 0.4, color = color_point) +
    geom_smooth(method = "loess", se = FALSE, color = color_line) +
    theme_minimal() +
    scale_x_continuous(breaks = seq(min(clean_data$testDate),
                                    max(clean_data$testDate), by = 1)) +
    labs(title = title, x = "Year", y = ylab)
  int_plot(p)
}
```

## Part 2: Loading in Datasets

(a) Load in Geekbench

```
recent_cpu <- read.csv(here("Datasets", "Geekbench", "recent-cpu-v6.csv"))
recent_gpu <- read.csv(here("Datasets", "Geekbench", "recent-gpu-v6.csv"))
single_core <- read.csv(here("Datasets", "Geekbench", "single-core-v4.csv"))
top_multi <- read.csv(here("Datasets", "Geekbench", "top-multi-core-v6.csv"))
top_single <- read.csv(here("Datasets", "Geekbench", "top-single-core-v6.csv"))
```

(b) Load in Kaggle

```
gpu_benchmarks <- read.csv(here("Datasets", "Kaggle", "GPU_benchmarks_v7.csv"))
gpu_scores <- read.csv(here("Datasets", "Kaggle", "GPU_scores_graphicsAPIs.csv"))
```

## Part 3: Preprocessing and Merging Datasets

(a) Make GPU names consistent across datasets (lowercase and trimmed). Merge the PassMark and Geekbench data based on the GPU name. After merging, removes duplicate name columns, prints out how many rows are in each dataset (before and after the merge), and shows a quick preview of the merged result.

```
gpu_benchmarks$gpu_name <- tolower(trimws(gpu_benchmarks$gpuName))
gpu_scores$gpu_name <- tolower(trimws(gpu_scores$Device))
merged_gpu <- merge(gpu_benchmarks, gpu_scores, by="gpu_name")
merged_gpu <- merged_gpu |> select(-gpuName, -Device)

cat("Rows in PassMark dataset:", nrow(gpu_benchmarks), "\n")

## Rows in PassMark dataset: 2317
cat("Rows in Geekbench dataset:", nrow(gpu_scores), "\n")

## Rows in Geekbench dataset: 1213
cat("Rows in merged dataset:", nrow(merged_gpu), "\n\n")
```

## Rows in merged dataset: 647

```
head(merged_gpu)
```

```
##          gpu_name G3Dmark G2Dmark price gpuValue TDP powerPerformance testDate
## 1      a40-12q    5573     198    NA      NA    NA      NA        NA    2022
## 2 firepro m4000    1597     410  72.83    21.92    NA      NA        NA    2012
## 3 firepro m4100    1059     623    NA      NA    NA      NA        NA    2015
## 4 firepro m4150     999     207    NA      NA    NA      NA        NA    2015
## 5 firepro m4170    1067     290    NA      NA    NA      NA        NA    2015
## 6 firepro m5100    2103     800    NA      NA    NA      NA        NA    2014
##          category Manufacturer CUDA Metal OpenCL Vulkan
## 1      Unknown       Nvidia  95329    NA 156643      NA
## 2 Workstation        AMD    NA    NA   6494      NA
## 3 Workstation        AMD    NA    NA   5067      NA
## 4      Unknown       AMD    NA    NA   5063   6685
## 5      Unknown       AMD    NA    NA   6347      NA
## 6 Workstation        AMD    NA    NA   9305  10692
```

## Part 4: Filtering the Datasets

- (a) Split the GPU data by manufacturer and generate a small summary showing how many rows and usable columns each manufacturer has.

```
manufacturers <- unique(gpu_scores$Manufacturer)
gpu_split <- split(gpu_scores, factor(gpu_scores$Manufacturer, levels = manufacturers))

for (m in manufacturers) {
  assign(sprintf("%s_gpu_scores", tolower(m)), subset(gpu_scores, Manufacturer == m))
}

manufacturers <- unique(gpu_scores$Manufacturer)
manufacturer_summary <- map_df(manufacturers, function(m) {
  df <- gpu_scores |> filter(Manufacturer == m)
  s <- df_stats(df)
  tibble(Manufacturer = m, Rows = s$rows, UsefulCols = s$useful_cols)
})

manufacturer_summary
```

## # A tibble: 9 x 3  
## Manufacturer Rows UsefulCols  
## <chr> <int> <int>  
## 1 Nvidia 404 7  
## 2 AMD 546 6  
## 3 Apple 21 5  
## 4 Qualcomm 22 4  
## 5 Intel 144 6  
## 6 Other 7 4  
## 7 ARM 58 5  
## 8 PowerVR 10 4  
## 9 Samsung 1 5

- (b) This splits the GPU benchmark data by supported API (CUDA, Metal, OpenCL, Vulkan) and summarizes each subset. For every test type, it keeps only GPUs with valid scores and reports how many useful rows and columns.

```
cuda_tests <- subset(gpu_scores, !is.na(CUDA))
metal_tests <- subset(gpu_scores, !is.na(Metal))
opencl_tests <- subset(gpu_scores, !is.na(OpenCL))
vulkan_tests <- subset(gpu_scores, !is.na(Vulkan))

test_types <- c("CUDA", "Metal", "OpenCL", "Vulkan")
test_summary <- map_df(test_types, function(t) {
  df <- gpu_scores |> filter(!is.na(.data[[t]]))
  if (nrow(df) == 0) return(NULL)
  s <- df_stats(df)
  tibble(Test = t, Rows = s$rows, UsefulCols = s$useful_cols)
})

test_summary
```

## # A tibble: 4 x 3  
## Test Rows UsefulCols  
## <chr> <int> <int>

```

## 1 CUDA      266      7
## 2 Metal     241      7
## 3 OpenCL    976      7
## 4 Vulkan    629      7

(c) This builds a summary table by manufacturer and benchmark type (CUDA, Metal, OpenCL, Vulkan).
For each valid combo, reports how many rows exist and how many columns useful.

summary_table <- map_df(manufacturers, function(m) {
  map_df(test_types, function(t) {
    df <- gpu_scores |> filter(Manufacturer == m, !is.na(.data[[t]]))
    if (nrow(df) == 0) return(NULL)
    s <- df_stats(df)
    tibble(Manufacturer = m, Test = t, Rows = s$rows, UsefulCols = s$useful_cols)
  })
})

summary_table

## # A tibble: 21 x 4
##   Manufacturer Test     Rows UsefulCols
##   <chr>        <chr>   <int>     <int>
## 1 Nvidia       CUDA     266       7
## 2 Nvidia       Metal     73        7
## 3 Nvidia       OpenCL    381       7
## 4 Nvidia       Vulkan    225       7
## 5 AMD          Metal     123       6
## 6 AMD          OpenCL    452       6
## 7 AMD          Vulkan    251       6
## 8 Apple         Metal     20        5
## 9 Apple         OpenCL    5         5
## 10 Qualcomm    OpenCL    1         4
## # i 11 more rows

```

## Part 5: Exploring Data Through ggplot() and plotly()

- (a) Comparing AMD vs Nvidia (CUDA/OpenCL/Vulkan/G3dmark)

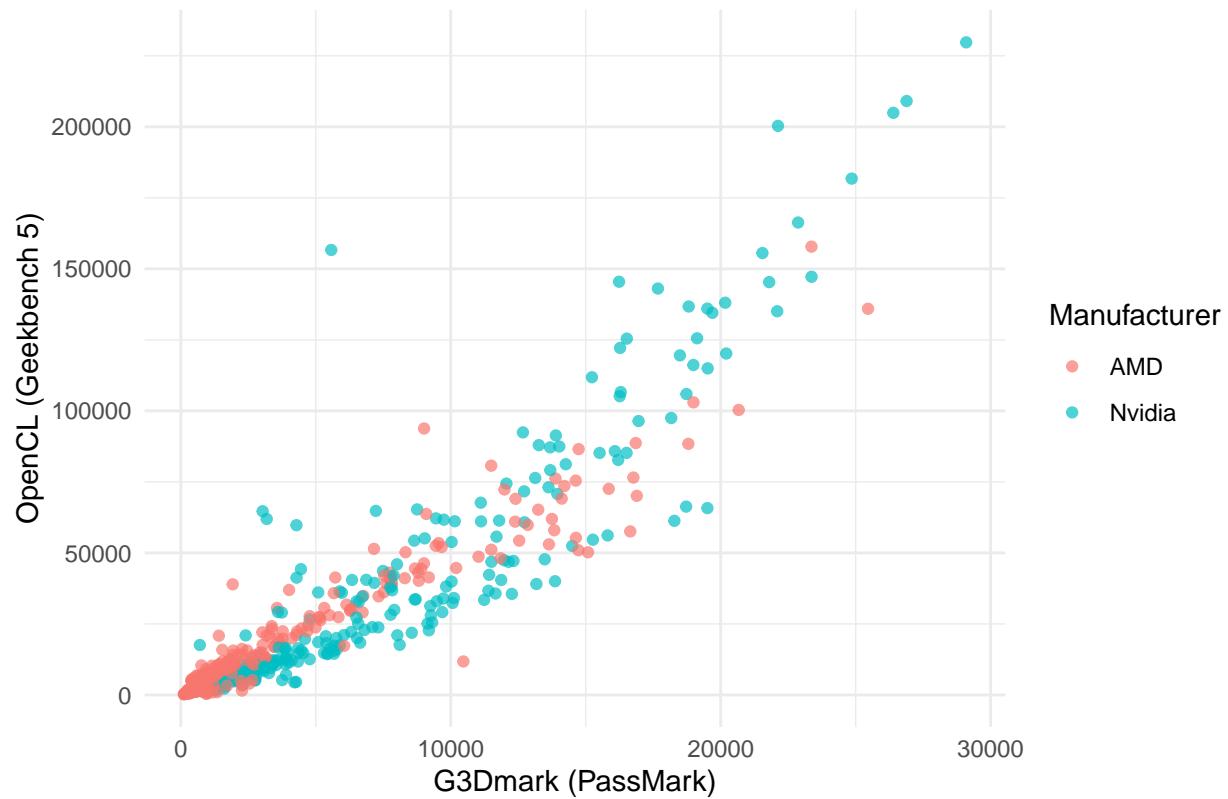
```

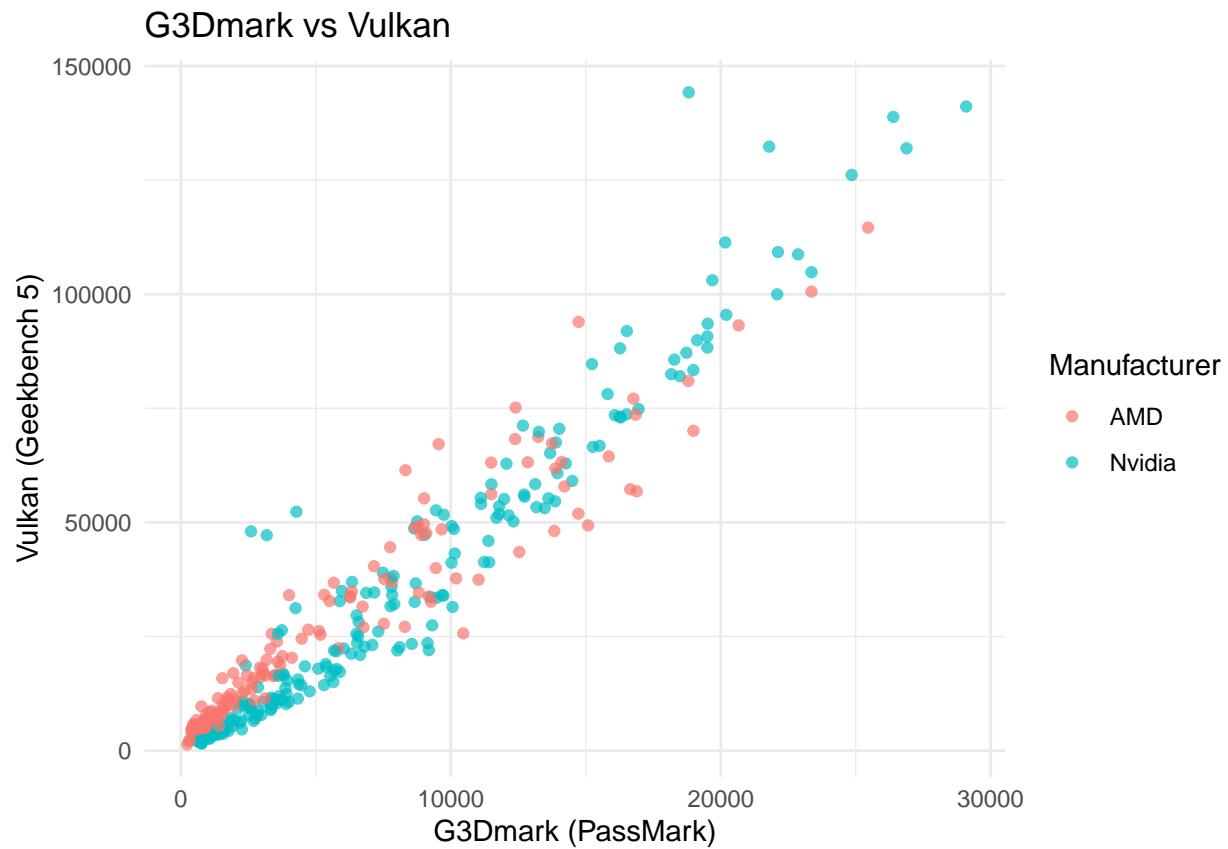
plot_scatter <- function(x_col, y_col, data) {
  p <- ggplot(data, aes_string(x = x_col, y = y_col, color = "Manufacturer")) +
    geom_point(alpha = 0.7) +
    theme_minimal() +
    labs(title = paste(x_col, "vs", y_col), x = paste(x_col, "(PassMark)"),
         y = paste(y_col, "(Geekbench 5)"))
  int_plot(p)
}

print(plot_scatter("G3Dmark", "OpenCL", merged_gpu))


```

### G3Dmark vs OpenCL





```
if (!is.null(merged_gpu$Metal)) {
  print(plot_scatter("G3Dmark", "Metal", merged_gpu))
}
```



(b) Plotting different trends over time

```
# add PerfPerWatt column once
merged_gpu <- merged_gpu |> mutate(PerfPerWatt = G3Dmark / TDP)

get_slope <- function(data, y, manufacturer = NULL, digits = 0) {
  d <- data |> filter(!is.na(testDate), !is.na(.data[[y]]))


  if (!is.null(manufacturer)) {
    d <- d |> filter(Manufacturer == manufacturer)
  }

  fit <- lm(reformulate("testDate", response = y), data = d)
  round(coef(fit)[["testDate"]], digits)
}

slope_g3d <- get_slope(merged_gpu, "G3Dmark")
slope_tdp <- get_slope(merged_gpu, "TDP", digits = 2)
slope_eff <- get_slope(merged_gpu, "PerfPerWatt", digits = 2)
slope_amd <- get_slope(merged_gpu, "G3Dmark", manufacturer = "AMD")
slope_nv <- get_slope(merged_gpu, "G3Dmark", manufacturer = "Nvidia")


# plots
plot_time_series(
  merged_gpu, "G3Dmark",
  paste("GPU Performance Trend (Avg: ", slope_g3d, " G3DMark/year)", ,
  "G3Dmark (PassMark Score)", "blue", "darkblue")
```

)




```
plot_time_series(  
  merged_gpu, "TDP",  
  paste("GPU Power Trend (Avg: ", slope_tdp, " W per year)"),  
  "TDP (Watts)", "red", "darkred"  
)
```

## GPU Power Trend (Avg: 9.02 W per year)




```
plot_time_series(  
  merged_gpu, "PerfPerWatt",  
  paste("GPU Efficiency Trend (Avg: ", slope_eff, " G3DMark per W/year)"),  
  "G3Dmark per Watt", "green", "darkgreen"  
)
```

## GPU Efficiency Trend (Avg: 9.51 G3DMark per W/year)



```
int_plot(
  ggplot(merged_gpu, aes(testDate, G3Dmark, color = Manufacturer)) +
  geom_point(alpha = 0.5) +
  geom_smooth(method = "loess", se = FALSE) +
  theme_minimal() +
  scale_x_continuous(breaks = seq(min(merged_gpu$testDate),
                                 max(merged_gpu$testDate), by = 1)) +
  labs(
    title = paste(
      "Performance by Manufacturer (AMD Avg: ",
      slope_amd, ", Nvidia Avg: ", slope_nv, " G3DMark/year)"),
    x = "Year", y = "G3Dmark Score")
)
```

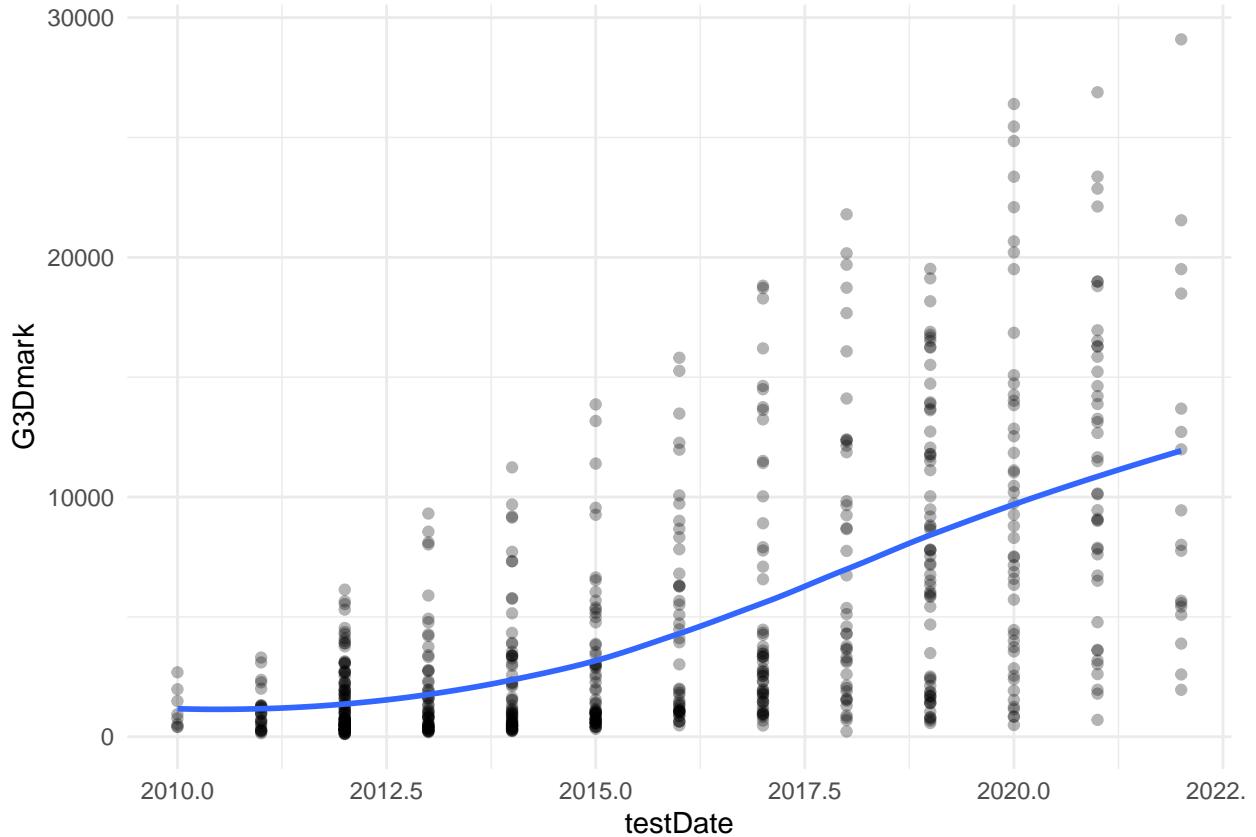
Performance by Manufacturer (AMD Avg: 900 , Nvidia Avg: 1141 G3DM



(c) Plotting ALL GPU Trend Line

```
merged_gpu <- merged_gpu |> mutate(PerfPerWatt = G3Dmark / TDP)

int_plot(
  ggplot(merged_gpu, aes(testDate, G3Dmark)) +
  geom_point(alpha=0.3) +
  geom_smooth(se=FALSE) +
  theme_minimal()
)
```



## Part 6: Exploring the data with TidyModels

- (a) Split/train/test on the merged dataset for our tidymodels. We clean up the merged dataset, keep only the variables we care about, drop missing values, and split the data into training and testing sets so we can fairly evaluate our models. Also, create a common recipe to use for all models (except Random Forest). Dummy variables are needed for linear models, but random forests already know how to work with categories.

```
set.seed(123)
model_data <- merged_gpu |>
  select(G3Dmark, testDate, TDP, price, Manufacturer, category) |> drop_na()
split <- initial_split(model_data, prop = 0.8)
train <- training(split)
test <- testing(split)

common_recipe <- recipe(G3Dmark ~ ., data = train) |>
  step_dummy(all_nominal_predictors()) |> step_normalize(all_numeric_predictors())
```

### Model 1: Linear Regression

- (b) We start with a simple linear regression as a baseline model, using dummy variables for categorical features and checking how well it predicts GPU performance on the test set.

```
lin_spec <- linear_reg() |> set_engine("lm")
lin_wf <- workflow() |> add_recipe(common_recipe) |> add_model(lin_spec)
lin_fit <- lin_wf |> fit(data = train)
```

```

lin_preds <- predict(lin_fit, new_data = test) |> bind_cols(test)
metrics(lin_preds, truth = G3Dmark, estimate = .pred)

## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard     2530.
## 2 rsq     standard      0.851
## 3 mae    standard     1927.

```

## Model 2, 3: LASSO/Ridge Regression

- (c) We try Ridge and LASSO regression, tuning the regularization strength with cross-validation to control overfitting and compare their predictive performance.

```

folds <- vfold_cv(train, v=10)
lambda_grid <- grid_regular(penalty(), levels= 30)

ridge_spec <- linear_reg(mixture=0, penalty =tune()) |> set_engine("glmnet")
lasso_spec <- linear_reg(mixture=1, penalty= tune()) |> set_engine("glmnet")

ridge_wf <- workflow() |> add_recipe(common_recipe) |> add_model(ridge_spec)
lasso_wf <- workflow() |> add_recipe(common_recipe) |> add_model(lasso_spec)

ridge_t <- tune_grid(ridge_wf, resamples=folds, grid=lambda_grid)
lasso_t <- tune_grid(lasso_wf, resamples=folds, grid=lambda_grid)

best_ridge <- select_best(ridge_t, metric="rmse")
best_lasso <- select_best(lasso_t, metric="rmse")

ridge_final_fit <- ridge_wf |> finalize_workflow(best_ridge) |> fit(data=train)
lasso_final_fit <- lasso_wf |> finalize_workflow(best_lasso) |> fit(data=train)

ridge_preds <- predict(ridge_final_fit, new_data=test) |> bind_cols(test)
lasso_preds <- predict(lasso_final_fit, new_data=test) |> bind_cols(test)

ridge_metrics <- ridge_preds |> metrics(truth=G3Dmark, estimate=.pred)
lasso_metrics <- lasso_preds |> metrics(truth=G3Dmark, estimate=.pred)

best_ridge

## # A tibble: 1 x 2
##   penalty .config
##   <dbl> <chr>
## 1 0.000000001 pre0_mod01_post0
best_lasso

## # A tibble: 1 x 2
##   penalty .config
##   <dbl> <chr>
## 1 0.000000001 pre0_mod01_post0
ridge_metrics

## # A tibble: 3 x 3

```

```

##   .metric .estimator .estimate
##   <chr>  <chr>      <dbl>
## 1 rmse    standard    2360.
## 2 rsq     standard     0.858
## 3 mae    standard    1751.

lasso_metrics

## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>  <chr>      <dbl>
## 1 rmse    standard    2522.
## 2 rsq     standard     0.850
## 3 mae    standard    1922.

```

## Model 4: KNN

- (d) We use KNN model, tuning the number of neighbors with cross-validation to see how a distance-based approach performs on this dataset.

```

knn_folds <- vfold_cv(train, v = 10)
knn_model <- nearest_neighbor(neighbors = tune()) |>
  set_engine("kknn") |> set_mode("regression")
knn_wf <- workflow() |> add_recipe(common_recipe) |> add_model(knn_model)

knn_tuned <- tune_grid(knn_wf, resamples = knn_folds,
                        grid = tibble(neighbors = c(1, 3, 5, 10, 20, 50)),
                        metrics = metric_set(rmse, rsq, mae))

show_best(knn_tuned, metric = "rmse")

## # A tibble: 5 x 7
##   neighbors .metric .estimator  mean    n std_err .config
##       <dbl>  <chr>  <chr>     <dbl> <int>  <dbl> <chr>
## 1        5 rmse    standard  2160.    10    270. pre0_mod3_post0
## 2        3 rmse    standard  2219.    10    254. pre0_mod2_post0
## 3       10 rmse    standard  2266.    10    262. pre0_mod4_post0
## 4       20 rmse    standard  2463.    10    252. pre0_mod5_post0
## 5        1 rmse    standard  2629.    10    223. pre0_mod1_post0

best_k <- select_best(knn_tuned, metric = "rmse")

knn_final <- finalize_workflow(knn_wf, best_k) |> fit(data = train)
knn_preds <- predict(knn_final, new_data = test) |> bind_cols(test)
metrics(knn_preds, truth = G3Dmark, estimate = .pred)

## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>  <chr>      <dbl>
## 1 rmse    standard    1709.
## 2 rsq     standard     0.918
## 3 mae    standard    1172.

```

## Model 5: Random Forest

- (e) Finally, we fit a random forest model to see which features matter most.

```

rf_recipe <- recipe(G3Dmark ~ ., data = train) |>
  step_normalize(all_numeric_predictors()) # we do not want dummy for rf.

rf_spec <- rand_forest(mtry = tune(), trees = tune(), min_n = tune()) |>
  set_engine("randomForest", importance = TRUE) |>
  set_mode("regression")

rf_wf <- workflow() |> add_recipe(rf_recipe) |> add_model(rf_spec)
folds <- vfold_cv(train, v = 10)

rf_params <- parameters(finalize(mtry(), train), trees(range = c(200, 600)),
  min_n(range = c(2, 25)))

rf_grid <- grid_regular(rf_params, levels = 4)

rf_tune <- tune_grid(rf_wf, resamples = folds, grid = rf_grid,
  metrics = metric_set(rmse, rsq, mae))

best_rf <- select_best(rf_tune, metric = "rmse")
final_rf <- finalize_workflow(rf_wf, best_rf)
rf_fit <- final_rf |> fit(data = train)
rf_preds <- predict(rf_fit, test) |> bind_cols(test)
rf_metrics <- metrics(rf_preds, truth = G3Dmark, estimate = .pred)
rf_metrics

## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard     1466.
## 2 rsq     standard      0.946
## 3 mae     standard     998.

rf_fit |> extract_fit_parsnip() |> pluck("fit") |> varImpPlot()

```

```
pluck(extract_fit_parsnip(rf_fit), "fit")
```

