

# STAT527 Term Project

Miller Kodish, Ian Ou, Vinay Pundith

12/17/2025

## Part 1: Define Helper Functions

- (a) These functions are used for tables. `useful_columns()` checks each column and keeps only the ones that don't have too many missing values, based on a threshold. `df_stats()` then uses that result to report simple summary info about the data, like how many rows it has and how many "useful" columns remain after filtering out columns with lots of NAs.

```
useful_columns <- function(df, na_threshold = 0.85) {
  na_fraction <- sapply(df, function(col) mean(is.na(col)))
  names(na_fraction[na_fraction <= na_threshold])
}

df_stats <- function(df, na_threshold=0.85) {
  cols <- useful_columns(df, na_threshold)
  list(rows = nrow(df), useful_cols = length(cols))
}
```

- (b) These functions are used for plotting. `int_plot()` automatically turns ggplot figures into interactive Plotly plots when the output is HTML, but keeps them static for PDFs. `plot_time_series()` cleans out missing values, makes a simple time-series plot with points and a smooth trend line so we don't have to repeat the same code.

```
# make sure plots are interactive in the R when compiled but static when in PDF
int_plot <- function(p) {
  if (knitr:::is_html_output()) {
    plotly::ggplotly(p)
  } else {
    p
  }
}

plot_time_series <- function(data, y, title, ylab, color_point, color_line) {
  clean_data <- data |> filter(!is.na(testDate), !is.na(.data[[y]]))
  p <- ggplot(clean_data, aes(x = testDate, y = .data[[y]])) +
    geom_point(alpha = 0.4, color = color_point) +
    geom_smooth(method = "loess", se = FALSE, color = color_line) +
    theme_minimal() +
    scale_x_continuous(breaks = seq(min(clean_data$testDate),
                                    max(clean_data$testDate), by = 1)) +
    labs(title = title, x = "Year", y = ylab)
  int_plot(p)
}
```

## Part 2: Loading in Datasets

(a) Load in Geekbench

```
recent_cpu <- read.csv(here("Datasets", "Geekbench", "recent-cpu-v6.csv"))
recent_gpu <- read.csv(here("Datasets", "Geekbench", "recent-gpu-v6.csv"))
single_core <- read.csv(here("Datasets", "Geekbench", "single-core-v4.csv"))
top_multi <- read.csv(here("Datasets", "Geekbench", "top-multi-core-v6.csv"))
top_single <- read.csv(here("Datasets", "Geekbench", "top-single-core-v6.csv"))
```

(b) Load in Kaggle

```
gpu_benchmarks <- read.csv(here("Datasets", "Kaggle", "GPU_benchmarks_v7.csv"))
gpu_scores <- read.csv(here("Datasets", "Kaggle", "GPU_scores_graphicsAPIs.csv"))
```

## Part 3: Preprocessing and Merging Datasets

(a) Make GPU names consistent across datasets (lowercase and trimmed). Merge the PassMark and Geekbench data based on the GPU name. After merging, removes duplicate name columns, prints out how many rows are in each dataset (before and after the merge), and shows a quick preview of the merged result.

```
gpu_benchmarks$gpu_name <- tolower(trimws(gpu_benchmarks$gpuName))
gpu_scores$gpu_name <- tolower(trimws(gpu_scores$Device))
merged_gpu <- merge(gpu_benchmarks, gpu_scores, by="gpu_name")
merged_gpu <- merged_gpu |> select(-gpuName, -Device)

cat("Rows in PassMark dataset:", nrow(gpu_benchmarks), "\n")

## Rows in PassMark dataset: 2317
cat("Rows in Geekbench dataset:", nrow(gpu_scores), "\n")

## Rows in Geekbench dataset: 1213
cat("Rows in merged dataset:", nrow(merged_gpu), "\n\n")

## Rows in merged dataset: 647
head(merged_gpu, 50)
```

```
##          gpu_name G3Dmark G2Dmark    price gpuValue TDP powerPerformance
## 1      a40-12q    5573     198       NA       NA   NA        NA
## 2  firepro m4000    1597     410    72.83    21.92   NA        NA
## 3  firepro m4100    1059     623       NA       NA   NA        NA
## 4  firepro m4150     999     207       NA       NA   NA        NA
## 5  firepro m4170    1067     290       NA       NA   NA        NA
## 6  firepro m5100    2103     800       NA       NA   NA        NA
## 7  firepro m5950    1314     574       NA       NA   NA        NA
## 8  firepro m6000    1820     776    325.99     5.58   NA        NA
## 9  firepro m6100    1945     338       NA       NA   NA        NA
## 10 firepro s10000    5315     493       NA       NA   NA        NA
## 11 firepro s7150    6276     866  1697.99     3.70   NA        NA
## 12 firepro v3900     662     242    238.44     2.78    50    13.24
## 13 firepro v4900     974     318   149.99     6.50    75    12.99
## 14 firepro v5800    1186     339    53.00    22.37   NA        NA
## 15 firepro v5900    1271     352    86.00    14.78    75    16.94
```

## 16	firepro v7800	2034	290	190.00	10.70	NA		NA
## 17	firepro w2100	890	327	59.00	15.09	26		34.24
## 18	firepro w4100	1541	535	184.00	8.38	50		30.82
## 19	firepro w4170m	953	223	NA	NA	NA		NA
## 20	firepro w4190m	1086	194	NA	NA	NA		NA
## 21	firepro w5000	3056	587	184.00	16.61	75		40.74
## 22	firepro w5100	2957	541	125.99	23.47	75		39.43
## 23	firepro w5130m	1341	307	NA	NA	NA		NA
## 24	firepro w5170m	1555	399	NA	NA	NA		NA
## 25	firepro w7000	4276	551	278.05	15.38	150		28.51
## 26	firepro w7100	5175	667	389.30	13.29	150		34.50
## 27	firepro w8100	7327	759	704.00	10.41	220		33.30
## 28	firepro w9000	6138	720	845.00	7.26	274		22.40
## 29	firepro w9100	7719	790	1729.99	4.46	275		28.07
## 30	geforce 410m	250	129	NA	NA	12		20.80
## 31	geforce 510	257	199	69.99	3.67	25		10.27
## 32	geforce 605	305	179	136.83	2.23	NA		NA
## 33	geforce 610m	292	103	NA	NA	NA		NA
## 34	geforce 615	552	412	NA	NA	NA		NA
## 35	geforce 710a	451	258	NA	NA	33		13.67
## 36	geforce 710m	450	116	NA	NA	12		37.48
## 37	geforce 730a	769	205	NA	NA	NA		NA
## 38	geforce 800m	468	388	NA	NA	NA		NA
## 39	geforce 810a	622	319	NA	NA	NA		NA
## 40	geforce 810m	402	71	NA	NA	NA		NA
## 41	geforce 820a	568	162	NA	NA	NA		NA
## 42	geforce 820m	511	111	NA	NA	15		34.04
## 43	geforce 830a	1005	736	NA	NA	NA		NA
## 44	geforce 830m	976	136	NA	NA	25		39.06
## 45	geforce 840a	1180	266	NA	NA	NA		NA
## 46	geforce 840m	1067	147	NA	NA	30		35.56
## 47	geforce 845m	1430	233	NA	NA	NA		NA
## 48	geforce 910m	578	150	NA	NA	NA		NA
## 49	geforce 920m	723	119	NA	NA	NA		NA
## 50	geforce 920mx	1088	149	NA	NA	NA		NA
##	testDate		category	Manufacturer	CUDA	Metal	OpenCL	Vulkan
## 1	2022		Unknown	Nvidia	95329	NA	156643	NA
## 2	2012		Workstation	AMD	NA	NA	6494	NA
## 3	2015		Workstation	AMD	NA	NA	5067	NA
## 4	2015		Unknown	AMD	NA	NA	5063	6685
## 5	2015		Unknown	AMD	NA	NA	6347	NA
## 6	2014		Workstation	AMD	NA	NA	9305	10692
## 7	2011		Workstation	AMD	NA	NA	1505	NA
## 8	2012		Workstation	AMD	NA	NA	9420	NA
## 9	2013		Workstation	AMD	NA	NA	15612	16951
## 10	2015		Unknown	AMD	NA	NA	30631	34145
## 11	2016		Unknown	AMD	NA	NA	29623	33575
## 12	2012		Workstation	AMD	NA	NA	1395	NA
## 13	2012		Workstation	AMD	NA	NA	1666	NA
## 14	2016		Unknown	AMD	NA	NA	2740	NA
## 15	2011		Workstation	AMD	NA	NA	2880	NA
## 16	2018		Unknown	AMD	NA	NA	4894	NA
## 17	2014		Workstation	AMD	NA	NA	4425	NA
## 18	2014		Workstation	AMD	NA	NA	6575	7852

## 19	2015	Unknown	AMD	NA	NA	5507	NA
## 20	2016	Unknown	AMD	NA	NA	5705	NA
## 21	2012	Workstation	AMD	NA	NA	12387	16955
## 22	2014	Workstation	AMD	NA	NA	13903	16106
## 23	2016	Unknown	AMD	NA	NA	7779	8337
## 24	2015	Mobile, Workstation	AMD	NA	NA	8775	NA
## 25	2012	Workstation	AMD	NA	NA	21082	NA
## 26	2015	Workstation	AMD	NA	NA	26357	25403
## 27	2014	Workstation	AMD	NA	NA	34705	NA
## 28	2012	Workstation	AMD	NA	NA	31775	NA
## 29	2014	Workstation	AMD	NA	NA	43046	NA
## 30	2011	Mobile	Nvidia	NA	NA	872	NA
## 31	2011	Desktop	Nvidia	NA	NA	758	NA
## 32	2012	Unknown	Nvidia	NA	NA	764	NA
## 33	2012	Unknown	Nvidia	NA	NA	1086	NA
## 34	2013	Unknown	Nvidia	NA	NA	1162	NA
## 35	2014	Desktop	Nvidia	NA	NA	2638	NA
## 36	2013	Mobile	Nvidia	NA	NA	1976	NA
## 37	2015	Unknown	Nvidia	NA	NA	3088	NA
## 38	2014	Unknown	Nvidia	NA	NA	1303	NA
## 39	2014	Unknown	Nvidia	NA	NA	1828	NA
## 40	2014	Unknown	Nvidia	NA	NA	1644	NA
## 41	2015	Unknown	Nvidia	NA	NA	2708	NA
## 42	2014	Mobile	Nvidia	NA	NA	2347	NA
## 43	2014	Unknown	Nvidia	NA	NA	4315	NA
## 44	2014	Mobile	Nvidia	4342	NA	4422	4078
## 45	2014	Unknown	Nvidia	NA	NA	5596	NA
## 46	2014	Mobile	Nvidia	5629	NA	5938	5207
## 47	2015	Mobile	Nvidia	NA	NA	6112	NA
## 48	2015	Unknown	Nvidia	NA	NA	2629	NA
## 49	2015	Unknown	Nvidia	2766	NA	3462	3215
## 50	2016	Unknown	Nvidia	4274	NA	4210	3723

## Part 4: Filtering the Datasets

- (a) Split the GPU data by manufacturer and generate a small summary showing how many rows and usable columns each manufacturer has.

```
manufacturers <- unique(gpu_scores$Manufacturer)
gpu_split <- split(gpu_scores, factor(gpu_scores$Manufacturer, levels = manufacturers))

for (m in manufacturers) {
  assign(sprintf("%s_gpu_scores", tolower(m)), subset(gpu_scores, Manufacturer == m))
}

manufacturers <- unique(gpu_scores$Manufacturer)
manufacturer_summary <- map_df(manufacturers, function(m) {
  df <- gpu_scores |> filter(Manufacturer == m)
  s <- df_stats(df)
  tibble(Manufacturer = m, Rows = s$rows, UsefulCols = s$useful_cols)
})

manufacturer_summary
```

```

## # A tibble: 9 x 3
##   Manufacturer  Rows UsefulCols
##   <chr>        <int>     <int>
## 1 Nvidia        404      7
## 2 AMD           546      6
## 3 Apple          21      5
## 4 Qualcomm       22      4
## 5 Intel          144      6
## 6 Other           7      4
## 7 ARM            58      5
## 8 PowerVR        10      4
## 9 Samsung         1      5

```

- (b) This splits the GPU benchmark data by supported API (CUDA, Metal, OpenCL, Vulkan) and summarizes each subset. For every test type, it keeps only GPUs with valid scores and reports how many useful rows and columns.

```

cuda_tests <- subset(gpu_scores, !is.na(CUDA))
metal_tests <- subset(gpu_scores, !is.na(Metal))
opencl_tests <- subset(gpu_scores, !is.na(OpenCL))
vulkan_tests <- subset(gpu_scores, !is.na(Vulkan))

test_types <- c("CUDA", "Metal", "OpenCL", "Vulkan")
test_summary <- map_df(test_types, function(t) {
  df <- gpu_scores |> filter(!is.na(.data[[t]]))
  if (nrow(df) == 0) return(NULL)
  s <- df_stats(df)
  tibble(Test = t, Rows = s$rows, UsefulCols = s$useful_cols)
})

test_summary

```

```

## # A tibble: 4 x 3
##   Test    Rows UsefulCols
##   <chr> <int>     <int>
## 1 CUDA    266      7
## 2 Metal    241      7
## 3 OpenCL   976      7
## 4 Vulkan   629      7

```

- (c) This builds a summary table by manufacturer and benchmark type (CUDA, Metal, OpenCL, Vulkan). For each valid combo, reports how many rows exist and how many columns useful.

```

summary_table <- map_df(manufacturers, function(m) {
  map_df(test_types, function(t) {
    df <- gpu_scores |> filter(Manufacturer == m, !is.na(.data[[t]]))
    if (nrow(df) == 0) return(NULL)
    s <- df_stats(df)
    tibble(Manufacturer = m, Test = t, Rows = s$rows, UsefulCols = s$useful_cols)
  })
})

summary_table

```

```

## # A tibble: 21 x 4
##   Manufacturer Test    Rows UsefulCols
##   <chr>        <chr> <int>     <int>
## 1 Nvidia        CUDA    266      7
## 2 Nvidia        Metal    241      7
## 3 Nvidia        OpenCL   976      7
## 4 Nvidia        Vulkan   629      7
## 5 AMD           CUDA    546      6
## 6 AMD           Metal    58      5
## 7 AMD           OpenCL   10      4
## 8 AMD           Vulkan   1      5
## 9 Apple          CUDA    21      5
## 10 Apple         Metal    7      4
## 11 Qualcomm       CUDA    22      4
## 12 Qualcomm       Metal    144      6
## 13 Qualcomm       OpenCL   1      5
## 14 Qualcomm       Vulkan   1      5
## 15 Intel          CUDA    144      6
## 16 Intel          Metal    58      5
## 17 Intel          OpenCL   1      4
## 18 Intel          Vulkan   1      5
## 19 Other           CUDA    7      4
## 20 Other           Metal    1      5
## 21 Other           OpenCL   1      4
## 22 Other           Vulkan   1      5

```

```

##      <chr>     <chr>   <int>   <int>
## 1 Nvidia      CUDA    266      7
## 2 Nvidia      Metal     73      7
## 3 Nvidia      OpenCL   381      7
## 4 Nvidia      Vulkan   225      7
## 5 AMD         Metal    123      6
## 6 AMD         OpenCL   452      6
## 7 AMD         Vulkan   251      6
## 8 Apple        Metal    20       5
## 9 Apple        OpenCL    5       5
## 10 Qualcomm   OpenCL    1       4
## # i 11 more rows

```

## Part 5: Exploring Data Through ggplot() and plotly()

- (a) Comparing AMD vs Nvidia (CUDA/OpenCL/Vulkan/G3dmark)

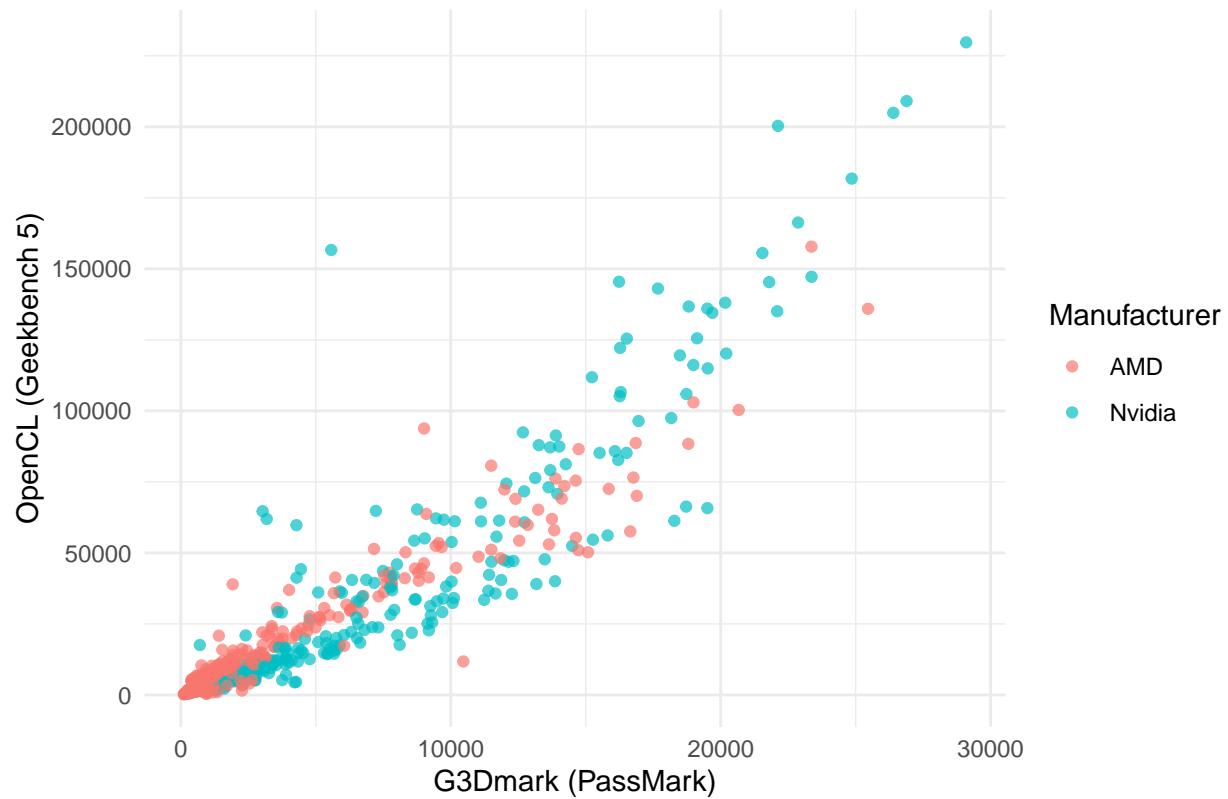
```

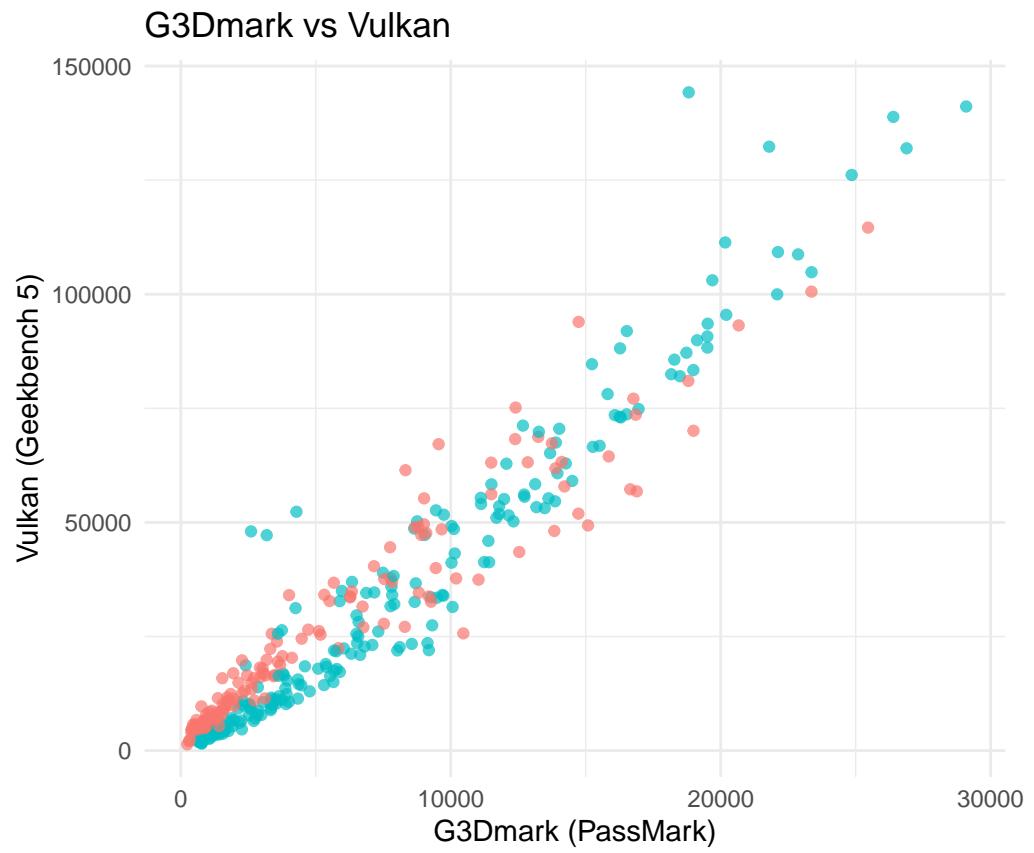
plot_scatter <- function(x_col, y_col, data) {
  p <- ggplot(data, aes_string(x = x_col, y = y_col, color = "Manufacturer")) +
    geom_point(alpha = 0.7) +
    theme_minimal() +
    labs(title = paste(x_col, "vs", y_col), x = paste(x_col, "(PassMark")),
        y = paste(y_col, "(Geekbench 5)"))
  int_plot(p)
}

print(plot_scatter("G3Dmark", "OpenCL", merged_gpu))

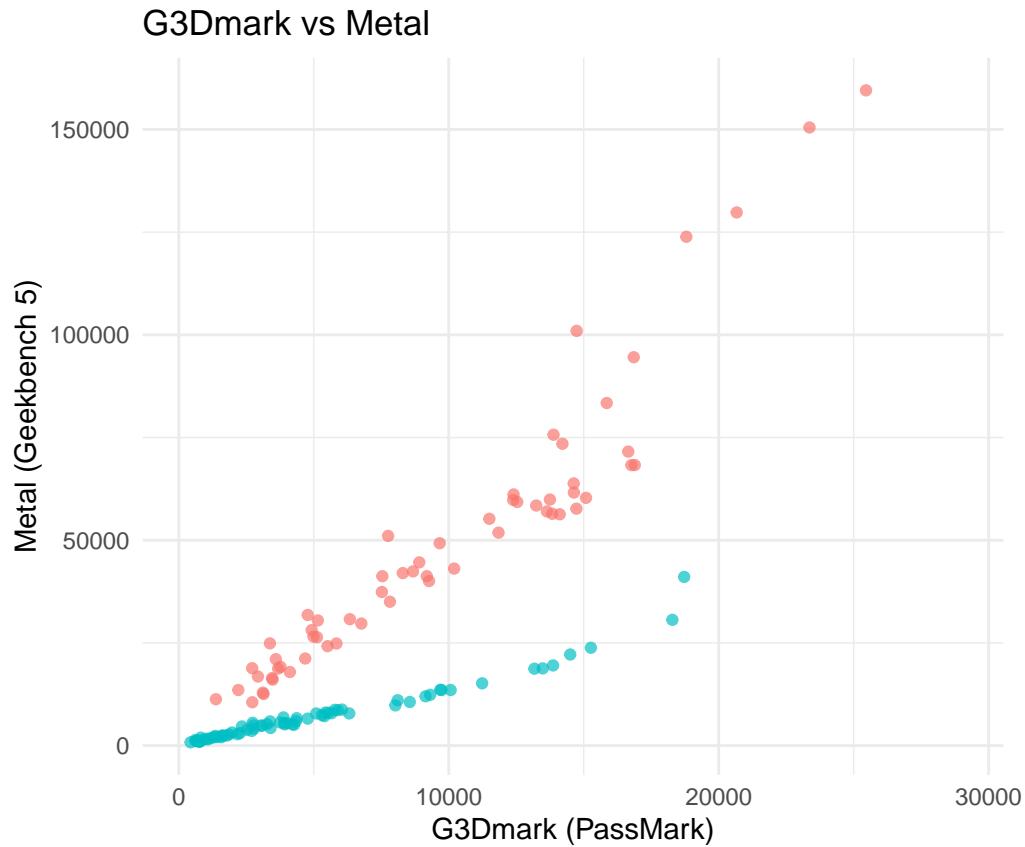
```

### G3Dmark vs OpenCL





```
if (!is.null(merged_gpu$Metal)) {  
  print(plot_scatter("G3Dmark", "Metal", merged_gpu))  
}
```



(b) Plotting different trends over time

```
# add PerfPerWatt column once
merged_gpu <- merged_gpu |> mutate(PerfPerWatt = G3Dmark / TDP)

get_slope <- function(data, y, manufacturer = NULL, digits = 0) {
  d <- data |> filter(!is.na(testDate), !is.na(.data[[y]]))

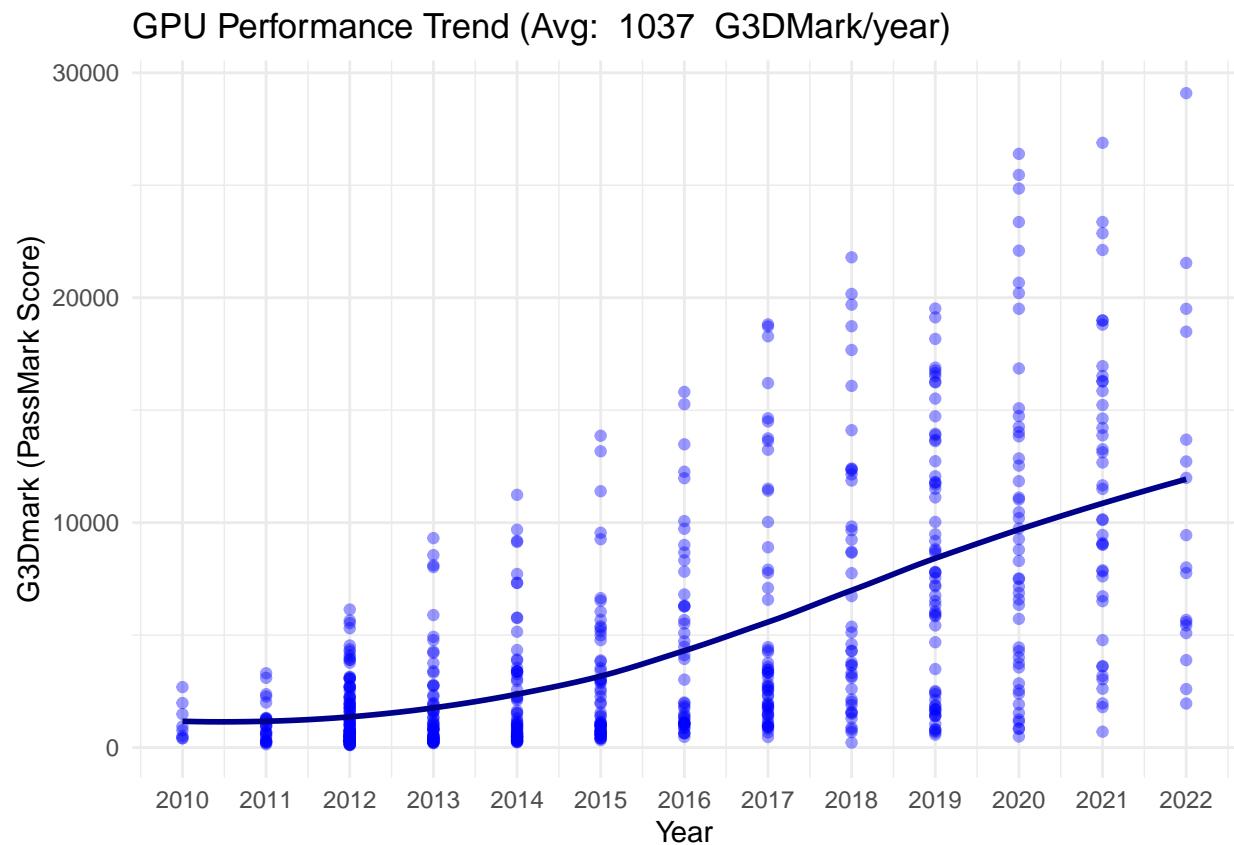
  if (!is.null(manufacturer)) {
    d <- d |> filter(Manufacturer == manufacturer)
  }

  fit <- lm(reformulate("testDate", response = y), data = d)
  round(coef(fit)[["testDate"]], digits)
}

slope_g3d <- get_slope(merged_gpu, "G3Dmark")
slope_tdp <- get_slope(merged_gpu, "TDP", digits = 2)
slope_eff <- get_slope(merged_gpu, "PerfPerWatt", digits = 2)
slope_amd <- get_slope(merged_gpu, "G3Dmark", manufacturer = "AMD")
slope_nv <- get_slope(merged_gpu, "G3Dmark", manufacturer = "Nvidia")

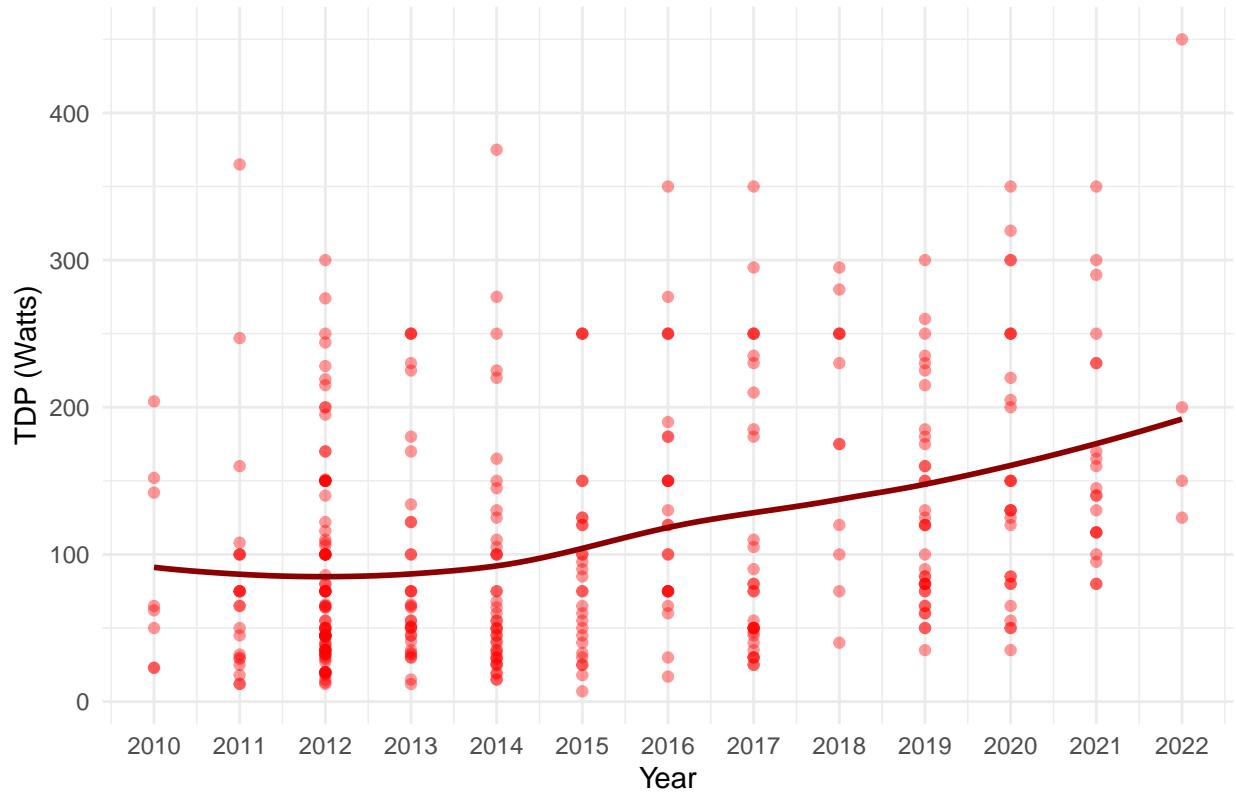
# plots
plot_time_series(
  merged_gpu, "G3Dmark",
  paste("GPU Performance Trend (Avg: ", slope_g3d, " G3DMark/year)", ,
  "G3Dmark (PassMark Score)", "blue", "darkblue")
```

)



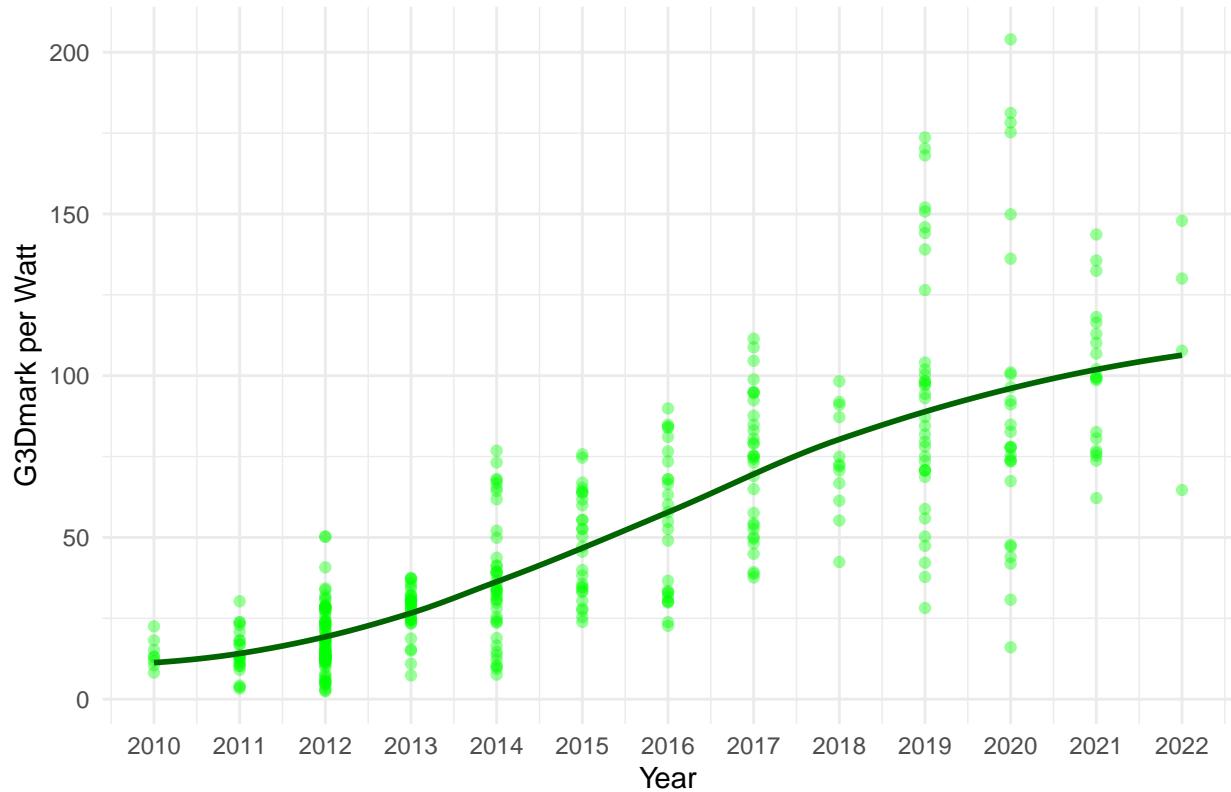
```
plot_time_series(  
  merged_gpu, "TDP",  
  paste("GPU Power Trend (Avg: ", slope_tdp, " W per year)"),  
  "TDP (Watts)", "red", "darkred"  
)
```

## GPU Power Trend (Avg: 9.02 W per year)

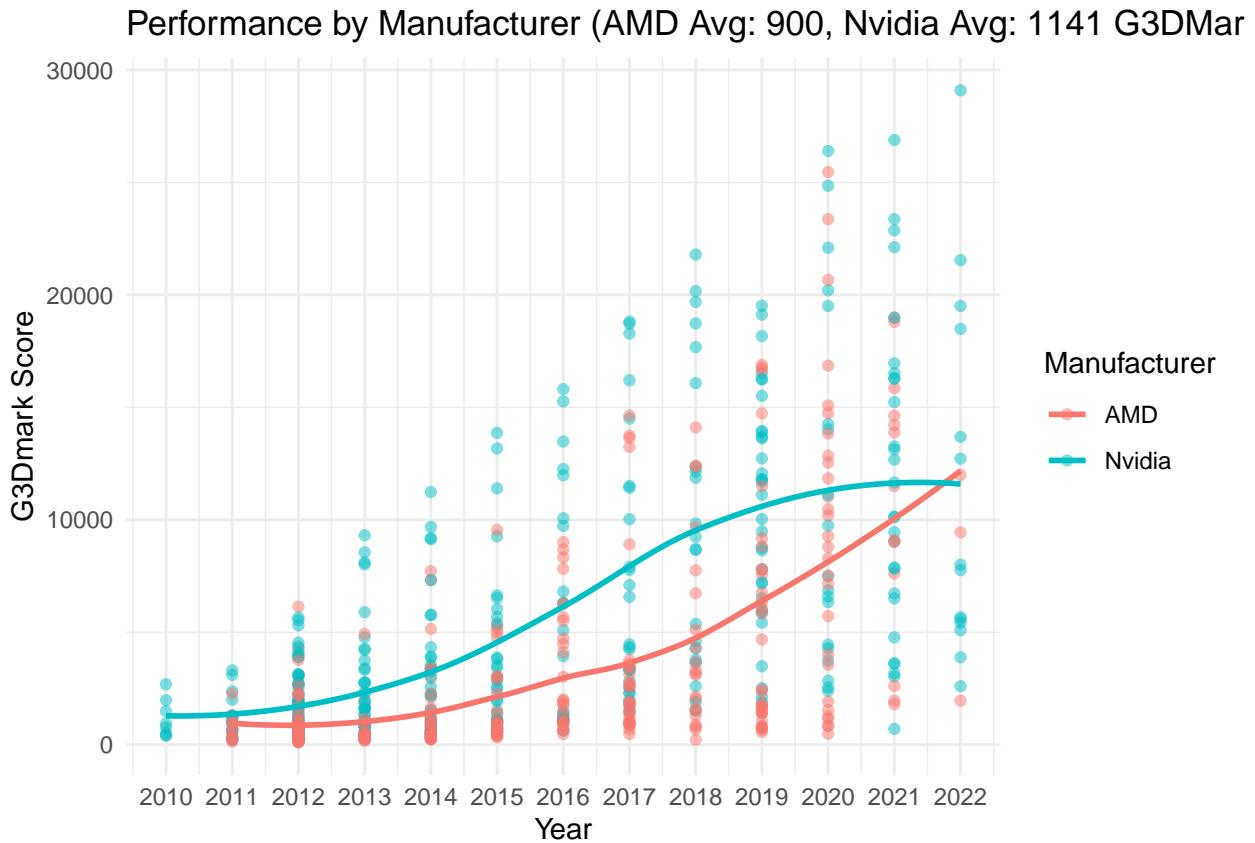


```
plot_time_series(  
  merged_gpu, "PerfPerWatt",  
  paste("GPU Efficiency Trend (Avg: ", slope_eff, " G3DMark per W/year)"),  
  "G3Dmark per Watt", "green", "darkgreen"  
)
```

## GPU Efficiency Trend (Avg: 9.51 G3DMark per W/year)



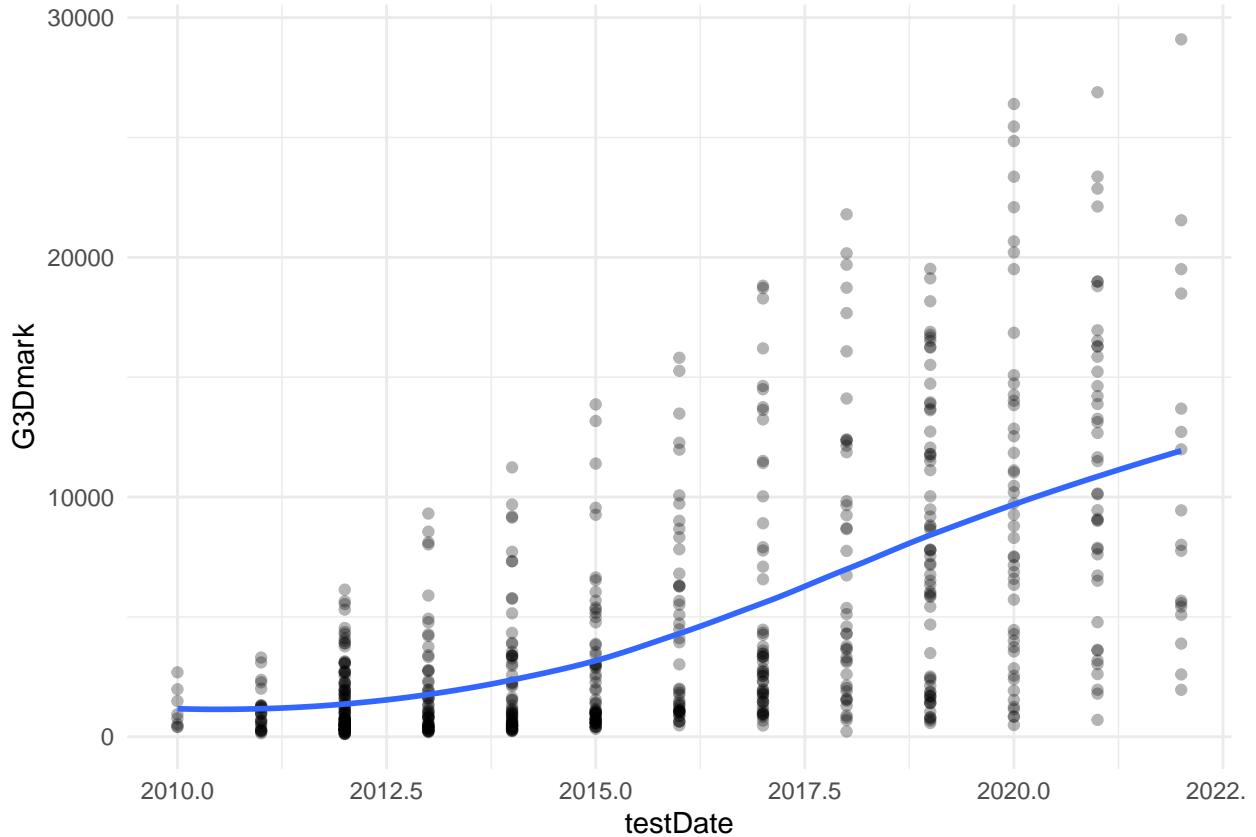
```
int_plot(
  ggplot(merged_gpu, aes(testDate, G3Dmark, color = Manufacturer)) +
  geom_point(alpha = 0.5) +
  geom_smooth(method = "loess", se = FALSE) +
  theme_minimal() +
  scale_x_continuous(breaks = seq(min(merged_gpu$testDate),
                                 max(merged_gpu$testDate), by = 1)) +
  labs(
    title = paste0(
      "Performance by Manufacturer (AMD Avg: ",
      slope_amd, ", Nvidia Avg: ", slope_nv, " G3DMark/year)"
    ),
    x = "Year", y = "G3Dmark Score"
  )
)
```



(c) Plotting ALL GPU Trend Line

```
merged_gpu <- merged_gpu |> mutate(PerfPerWatt = G3Dmark / TDP)

int_plot(
  ggplot(merged_gpu, aes(testDate, G3Dmark)) +
  geom_point(alpha=0.3) +
  geom_smooth(se=FALSE) +
  theme_minimal()
)
```



## Part 6: Exploring the data with TidyModels

- (a) Split/train/test on the merged dataset for our tidymodels. We clean up the merged dataset, keep only the variables we care about, drop missing values, and split the data into training and testing sets so we can fairly evaluate our models. Also, create a common recipe to use for all models (except Random Forest). Dummy variables are needed for linear models, but random forests already know how to work with categories.

```
set.seed(123)
model_data <- merged_gpu |>
  select(G3Dmark, testDate, TDP, price, Manufacturer, category) |> drop_na()
split <- initial_split(model_data, prop = 0.8)
train <- training(split)
test <- testing(split)

common_recipe <- recipe(G3Dmark ~ ., data = train) |>
  step_dummy(all_nominal_predictors()) |> step_normalize(all_numeric_predictors())
```

### Model 1: Linear Regression

- (b) We start with a simple linear regression as a baseline model, using dummy variables for categorical features and checking how well it predicts GPU performance on the test set.

```
lin_spec <- linear_reg() |> set_engine("lm")
lin_wf <- workflow() |> add_recipe(common_recipe) |> add_model(lin_spec)
lin_fit <- lin_wf |> fit(data = train)
```

```

lin_preds <- predict(lin_fit, new_data = test) |> bind_cols(test)
metrics(lin_preds, truth = G3Dmark, estimate = .pred)

## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard     2530.
## 2 rsq     standard      0.851
## 3 mae    standard     1927.

```

### Model 3, 4: LASSO/Ridge Regression

- (d) We try Ridge and LASSO regression, tuning the regularization strength with cross-validation to control overfitting and compare their predictive performance.

```

folds <- vfold_cv(train, v=10)
lambda_grid <- grid_regular(penalty(), levels= 30)

ridge_spec <- linear_reg(mixture=0, penalty =tune()) |> set_engine("glmnet")
lasso_spec <- linear_reg(mixture=1, penalty= tune()) |> set_engine("glmnet")

ridge_wf <- workflow() |> add_recipe(common_recipe) |> add_model(ridge_spec)
lasso_wf <- workflow() |> add_recipe(common_recipe) |> add_model(lasso_spec)

ridge_t <- tune_grid(ridge_wf, resamples=folds, grid=lambda_grid)
lasso_t <- tune_grid(lasso_wf, resamples=folds, grid=lambda_grid)

best_ridge <- select_best(ridge_t, metric="rmse")
best_lasso <- select_best(lasso_t, metric="rmse")

ridge_final_fit <- ridge_wf |> finalize_workflow(best_ridge) |> fit(data=train)
lasso_final_fit <- lasso_wf |> finalize_workflow(best_lasso) |> fit(data=train)

ridge_preds <- predict(ridge_final_fit, new_data=test) |> bind_cols(test)
lasso_preds <- predict(lasso_final_fit, new_data=test) |> bind_cols(test)

ridge_metrics <- ridge_preds |> metrics(truth=G3Dmark, estimate=.pred)
lasso_metrics <- lasso_preds |> metrics(truth=G3Dmark, estimate=.pred)

best_ridge

## # A tibble: 1 x 2
##       penalty .config
##       <dbl> <chr>
## 1 0.000000001 pre0_mod01_post0
best_lasso

## # A tibble: 1 x 2
##       penalty .config
##       <dbl> <chr>
## 1 0.000000001 pre0_mod01_post0
ridge_metrics

## # A tibble: 3 x 3

```

```

##   .metric .estimator .estimate
##   <chr>  <chr>      <dbl>
## 1 rmse    standard    2360.
## 2 rsq     standard     0.858
## 3 mae    standard    1751.

lasso_metrics

## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>  <chr>      <dbl>
## 1 rmse    standard    2522.
## 2 rsq     standard     0.850
## 3 mae    standard    1922.

```

## Model 5: KNN

- (e) We use KNN model, tuning the number of neighbors with cross-validation to see how a distance-based approach performs on this dataset.

```

knn_folds <- vfold_cv(train, v = 10)
knn_model <- nearest_neighbor(neighbors = tune()) |>
  set_engine("kknn") |> set_mode("regression")
knn_wf <- workflow() |> add_recipe(common_recipe) |> add_model(knn_model)

knn_tuned <- tune_grid(knn_wf, resamples = knn_folds,
                        grid = tibble(neighbors = c(1, 3, 5, 10, 20, 50)),
                        metrics = metric_set(rmse, rsq, mae))

show_best(knn_tuned, metric = "rmse")

## # A tibble: 5 x 7
##   neighbors .metric .estimator  mean    n std_err .config
##       <dbl>  <chr>  <chr>     <dbl> <int>  <dbl> <chr>
## 1        5 rmse    standard  2160.    10    270. pre0_mod3_post0
## 2        3 rmse    standard  2219.    10    254. pre0_mod2_post0
## 3       10 rmse    standard  2266.    10    262. pre0_mod4_post0
## 4       20 rmse    standard  2463.    10    252. pre0_mod5_post0
## 5        1 rmse    standard  2629.    10    223. pre0_mod1_post0

best_k <- select_best(knn_tuned, metric = "rmse")

knn_final <- finalize_workflow(knn_wf, best_k) |> fit(data = train)
knn_preds <- predict(knn_final, new_data = test) |> bind_cols(test)
metrics(knn_preds, truth = G3Dmark, estimate = .pred)

## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>  <chr>      <dbl>
## 1 rmse    standard    1709.
## 2 rsq     standard     0.918
## 3 mae    standard    1172.

```

## Model 2: Random Forest

- (c) Finally, we fit a random forest model to see which features matter most.

```

rf_recipe <- recipe(G3Dmark ~ ., data = train) |>
  step_normalize(all_numeric_predictors()) # we do not want dummy for rf.

rf_spec <- rand_forest(mtry = tune(), trees = tune(), min_n = tune()) |>
  set_engine("randomForest", importance = TRUE) |>
  set_mode("regression")

rf_wf <- workflow() |> add_recipe(rf_recipe) |> add_model(rf_spec)
folds <- vfold_cv(train, v = 10)

rf_params <- parameters(finalize(mtry(), train), trees(range = c(200, 600)),
  min_n(range = c(2, 25)))

rf_grid <- grid_regular(rf_params, levels = 4)

rf_tune <- tune_grid(rf_wf, resamples = folds, grid = rf_grid,
  metrics = metric_set(rmse, rsq, mae))

best_rf <- select_best(rf_tune, metric = "rmse")
final_rf <- finalize_workflow(rf_wf, best_rf)
rf_fit <- final_rf |> fit(data = train)
rf_preds <- predict(rf_fit, test) |> bind_cols(test)
rf_metrics <- metrics(rf_preds, truth = G3Dmark, estimate = .pred)
rf_metrics

## # A tibble: 3 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 rmse    standard     1466.
## 2 rsq     standard      0.946
## 3 mae     standard     998.

rf_fit |> extract_fit_parsnip() |> pluck("fit") |> varImpPlot()

```

```
pluck(extract_fit_parsnip(rf_fit), "fit")
```

